


```

Epoch 1/20
Recall: 0.5250 - val_loss: 0.9030 - val_accuracy: 0.9374 - val_precision: 0.7497 - val_recall: 0.7497
Epoch 2/20
Recall: 0.5850 - val_loss: 0.8930 - val_accuracy: 0.9409 - val_precision: 0.8349 - val_recall: 0.7333
Epoch 3/20
Recall: 0.6700 - val_loss: 0.8694 - val_accuracy: 0.9656 - val_precision: 0.9656 - val_recall: 0.7849
Epoch 4/20
Recall: 0.8482 - val_loss: 0.8439 - val_accuracy: 0.9841 - val_precision: 0.9597 - val_recall: 0.9111
Epoch 5/20
Recall: 0.9095 - val_loss: 0.8206 - val_accuracy: 0.9862 - val_precision: 0.9526 - val_recall: 0.9359
Epoch 6/20
Recall: 0.9253 - val_loss: 0.1697 - val_accuracy: 0.9907 - val_precision: 0.9735 - val_recall: 0.9514
Epoch 7/20
Recall: 0.9428 - val_loss: 0.1558 - val_accuracy: 0.9911 - val_precision: 0.9706 - val_recall: 0.9576
Epoch 8/20
Recall: 0.9472 - val_loss: 0.1870 - val_accuracy: 0.9875 - val_precision: 0.9598 - val_recall: 0.9390
Epoch 9/20
Recall: 0.9535 - val_loss: 0.2177 - val_accuracy: 0.9840 - val_precision: 0.9752 - val_recall: 0.8945
Epoch 10/20
Recall: 0.9545 - val_loss: 0.1068 - val_accuracy: 0.9926 - val_precision: 0.9800 - val_recall: 0.9607
Epoch 11/20
Recall: 0.9685 - val_loss: 0.0974 - val_accuracy: 0.9931 - val_precision: 0.9760 - val_recall: 0.9690
Epoch 12/20
Recall: 0.9708 - val_loss: 0.1058 - val_accuracy: 0.9925 - val_precision: 0.9739 - val_recall: 0.9659
Epoch 13/20
Recall: 0.9664 - val_loss: 0.1127 - val_accuracy: 0.9919 - val_precision: 0.9718 - val_recall: 0.9628
Epoch 14/20
Recall: 0.9690 - val_loss: 0.1002 - val_accuracy: 0.9930 - val_precision: 0.9750 - val_recall: 0.9690
Epoch 15/20
Recall: 0.9767 - val_loss: 0.0788 - val_accuracy: 0.9944 - val_precision: 0.9793 - val_recall: 0.9762
Epoch 16/20
Recall: 0.9809 - val_loss: 0.0802 - val_accuracy: 0.9955 - val_precision: 0.9844 - val_recall: 0.9793
Epoch 17/20
Recall: 0.9829 - val_loss: 0.1071 - val_accuracy: 0.9925 - val_precision: 0.9729 - val_recall: 0.9669
Epoch 18/20
Recall: 0.9858 - val_loss: 0.0726 - val_accuracy: 0.9963 - val_precision: 0.9865 - val_recall: 0.9835
Epoch 19/20
Recall: 0.9871 - val_loss: 0.1026 - val_accuracy: 0.9944 - val_precision: 0.9793 - val_recall: 0.9762
Epoch 20/20
Recall: 0.9814 - val_loss: 0.0860 - val_accuracy: 0.9956 - val_precision: 0.9844 - val_recall: 0.9804
25/25
Test loss, Test Accuracy, Test Precision, Test Recall: [0.07072429358959198, 0.9978125955367432, 0.991249789423279, 0.991249789423279]

```

Evaluate Model

Evaluate model with different metrics:

- Accuracy
- Precision
- Recall
- F1-score

```

d_l_pipeline.evaluate_model(bilstm_attention_model, X_test, y_test)

```

```

13/13 [=====] - 5s 266ms/step
precision    recall    f1-score   support

0           0.97         1.00         0.99         33
1           1.00         0.89         0.94          9
2           0.98         0.98         0.98         48
3           1.00         1.00         1.00         38
4           1.00         0.99         1.00         632
5           1.00         1.00         1.00         1
6           1.00         0.97         0.99         36
7           0.43         1.00         0.60          3

accuracy          0.99         0.99         800
macro avg         0.92         0.98         0.94         800
weighted avg      0.99         0.99         0.99         800

```

Confusion Matrix

	0	1	2	3	4	5	6	7
0	33	0	0	0	0	0	0	0
1	0	8	0	0	0	0	0	0
2	0	0	47	0	1	0	0	0
3	0	0	0	38	0	0	0	0
4	0	0	0	0	628	0	0	4
5	0	0	0	0	0	1	0	0
6	0	0	1	0	0	0	35	0
7	0	0	0	0	0	0	0	3


Plot Accuracy and Loss

```
dl_pipeline.plot
```

1.00

accuracy

epoch	model loss (blue)	model loss (orange)
0.0	0.920	0.938
1.0	0.945	0.950
2.0	0.960	0.955



```
w2vec_embedding_dict
```

```
tive300.bin
Loading word2vec model
2000 00 00 00 11 56 310  INT32  1  1  1
```

Creating Embedding Matrix

We will create an embedding matrix from the Word2Vec vectors that will be used as weights in the embedding layer of our neural networks.

```
embedding_dim = 300
embedding_matrix = dl_pipeline.create_embedding_matrix(tokenizer, w2vec_embedding_dict, vocab_size, embedding_dim)
```

BILSTM

Let us create an BiLSTM model with the embeddings from Fasttext.

```
embedding_dim = 300
lstm_units = 128
dense = 8
w2vec_bilstm_model = dl_pipeline.create_BiLSTM_nn_model(embedding_matrix, vocab_size, max_len, embedding_dim,
                                                         lstm_units, dense)
w2vec_bilstm_model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
embedding_6 (Embedding)	(None, 100, 300)	232000
bidirectional_6 (Bidirectional)	(None, 256)	439296

```
dense_6 (Dense)          (None, 8)
```

```
=====
Total params: 670,552
Trainable params: 441,352
Non-trainable params: 229,200
=====

epochs = 20
batch_size = 128
history = di_pipeline.train_evaluate_model(w2vec_bilstm_model, X_train, y_train, X_test,
                                           y_test, epochs, batch_size)

Epoch 1/20
31/20 [=====] - 34s 914ms/step - loss: 1.0152 - accuracy: 0.9310 - precision: 0.8196 - recall: 0.5743 - val_loss: 0.6508 - val_accuracy: 0.9494 - val_precision: 0.8287 - val_recall: 0.7404
Epoch 2/20
31/20 [=====] - 32s 1s/step - loss: 0.5153 - accuracy: 0.9646 - precision: 0.9060 - recall: 0.8001 - val_loss: 0.3357 - val_accuracy: 0.9761 - val_precision: 0.9463 - val_recall: 0.8573
```

```
Epoch 3/20
31/31 [=====]
recall: 0.8774 - val_loss:
Epoch 4/20
```

```
31/31 [100%] 29s 915ms step - loss: 0.2083 accuracy: 0.9845 precision: 0.9603 - recall: 0.9141 - val_loss: 0.1475 - val_accuracy: 0.9886 - val_precision: 0.9661 - val_recall: 0.9421
31/31 [100%] 29s 915ms step - loss: 0.2083 accuracy: 0.9845 precision: 0.9603 - recall: 0.9141 - val_loss: 0.1475 - val_accuracy: 0.9886 - val_precision: 0.9661 - val_recall: 0.9421
31/31 [100%] 28s 806ms step - loss: 0.1449 accuracy: 0.9895 precision: 0.9683 - recall: 0.9467 - val_loss: 0.1127 - val_accuracy: 0.9926 - val_precision: 0.9789 - val_recall: 0.9617
31/31 [100%] 28s 806ms step - loss: 0.1449 accuracy: 0.9895 precision: 0.9683 - recall: 0.9467 - val_loss: 0.1127 - val_accuracy: 0.9926 - val_precision: 0.9789 - val_recall: 0.9617
31/31 [100%] 28s 908ms step - loss: 0.1037 accuracy: 0.9931 precision: 0.9778 - recall: 0.9666 - val_loss: 0.0846 - val_accuracy: 0.9934 - val_precision: 0.9781 - val_recall: 0.9690
31/31 [100%] 28s 908ms step - loss: 0.1037 accuracy: 0.9931 precision: 0.9778 - recall: 0.9666 - val_loss: 0.0846 - val_accuracy: 0.9934 - val_precision: 0.9781 - val_recall: 0.9690
31/31 [100%] 30s 967ms step - loss: 0.0762 accuracy: 0.9951 precision: 0.9836 - recall: 0.9772 - val_loss: 0.0752 - val_accuracy: 0.9950 - val_precision: 0.9833 - val_recall: 0.9762
31/31 [100%] 30s 967ms step - loss: 0.0762 accuracy: 0.9951 precision: 0.9836 - recall: 0.9772 - val_loss: 0.0752 - val_accuracy: 0.9950 - val_precision: 0.9833 - val_recall: 0.9762
31/31 [100%] 27s 882ms step - loss: 0.0571 accuracy: 0.9964 precision: 0.9875 - recall: 0.9834 - val_loss: 0.0661 - val_accuracy: 0.9953 - val_precision: 0.9844 - val_recall: 0.9783
31/31 [100%] 27s 882ms step - loss: 0.0571 accuracy: 0.9964 precision: 0.9875 - recall: 0.9834 - val_loss: 0.0661 - val_accuracy: 0.9953 - val_precision: 0.9844 - val_recall: 0.9783
31/31 [100%] 28s 910ms step - loss: 0.0503 accuracy: 0.9969 precision: 0.9889 - recall: 0.9866 - val_loss: 0.0621 - val_accuracy: 0.9957 - val_precision: 0.9844 - val_recall: 0.9814
31/31 [100%] 28s 910ms step - loss: 0.0503 accuracy: 0.9969 precision: 0.9889 - recall: 0.9866 - val_loss: 0.0621 - val_accuracy: 0.9957 - val_precision: 0.9844 - val_recall: 0.9814
```

```

31/31 [=====] 27s 88ms/step - loss: 0.0383 - accuracy: 0.9977 - precision: 0.9920 - recall: 0.9899 - val_loss: 0.0535 - val_accuracy: 0.9964 - val_precision: 0.9986 - val_recall: 0.9824
Epoch 11/20
31/31 [=====] 27s 87ms/step - loss: 0.0336 - accuracy: 0.9981 - precision: 0.9930 - recall: 0.9915 - val_loss: 0.0535 - val_accuracy: 0.9960 - val_precision: 0.9965 - val_recall: 0.9814
Epoch 12/20
31/31 [=====] 28s 916ms/step - loss: 0.0283 - accuracy: 0.9982 - precision: 0.9933 - recall: 0.9922 - val_loss: 0.0481 - val_accuracy: 0.9963 - val_precision: 0.9975 - val_recall: 0.9824
Epoch 13/20
31/31 [=====] 28s 904ms/step - loss: 0.0241 - accuracy: 0.9986 - precision: 0.9946 - recall: 0.9943 - val_loss: 0.0487 - val_accuracy: 0.9966 - val_precision: 0.9976 - val_recall: 0.9855
Epoch 14/20
31/31 [=====] 29s 949ms/step - loss: 0.0206 - accuracy: 0.9989 - precision: 0.9956 - recall: 0.9950 - val_loss: 0.0512 - val_accuracy: 0.9959 - val_precision: 0.9945 - val_recall: 0.9824
Epoch 15/20
31/31 [=====] 28s 918ms/step - loss: 0.0169 - accuracy: 0.9991 - precision: 0.9966 - recall: 0.9961 - val_loss: 0.0481 - val_accuracy: 0.9966 - val_precision: 0.9976 - val_recall: 0.9855
Epoch 16/20
31/31 [=====] 27s 879ms/step - loss: 0.0146 - accuracy: 0.9995 - precision: 0.9972 - recall: 0.9974 - val_loss: 0.0443 - val_accuracy: 0.9966 - val_precision: 0.9986 - val_recall: 0.9845
Epoch 17/20
24/31 [=====].....) - ETA: 5s - loss: 0.0134 - accuracy: 0.9995 - precision: 0.9980 - recall 1: 0.9977

```

```

Evaluate model with different metrics:

• Accuracy
• Precision
• Recall
• F1-score

d1_pipeline.evaluate_model(w2vec_bilstm_model, x_test, y_test)

13/13 [=====] - 8s 390ms/step

    precision    recall  f1-score   support

0         1.00        1.00        1.00        33
1         1.00        1.00        1.00        9
2         1.00        1.00        1.00        48
3         1.00        0.95        0.97        38
4         1.00        0.99        1.00        632
5         1.00        1.00        1.00        1
6         1.00        1.00        1.00        36
7         0.43        1.00        0.60        3

accuracy          0.99          0.99          0.99          800
macro avg         0.93          0.99          0.95          800
weighted avg      1.00          0.99          0.99          800

```

0	33	0	0
---	----	---	---

1	0	9	0	0	0	0	0
2	0	0	48	0	0	0	0

The confusion matrix shows the relationship between predicted and true labels. The true labels are on the y-axis (0-7) and predicted labels are on the x-axis (0-7). The color scale represents the count of instances, ranging from 0 to 400. The matrix is mostly diagonal, indicating high accuracy, with some off-diagonal elements, notably a cluster of values between true labels 3 and 4.

	0	1	2	3	4	5	6	7
0	0	0	0	36	2	0	0	0
1	0	0	0	0	0	0	0	4
2	0	0	0	0	360	0	0	0
3	0	0	0	0	0	0	1	0
4	0	0	0	0	0	0	36	0
5	0	0	0	0	0	0	0	3
6	0	0	0	0	0	0	0	0
7	0	0	0	0	0	0	0	0

Below the confusion matrix, a plot titled "model accuracy" shows the accuracy of the model over 10 epochs. The x-axis is labeled "epoch" and the y-axis is labeled "accuracy". The accuracy starts at approximately 0.985 and increases to about 0.995 by epoch 10. The plot shows a slight dip in accuracy around epoch 8, followed by a recovery.

The figure consists of two vertically stacked line plots sharing a common x-axis representing 'epoch' from 0 to 20.

The top plot shows 'accuracy' on the y-axis, ranging from 0.93 to 0.98. It contains two lines: an orange line and a blue line. Both lines show an upward trend, with the orange line generally performing better than the blue line.

Epoch	Orange Line Accuracy	Blue Line Accuracy
0	0.948	0.931
1	0.978	0.964
2	0.982	0.974
3	0.984	0.980

The bottom plot shows 'model loss' on the y-axis, ranging from 0.0 to 1.0. It contains two lines: an orange line and a blue line. Both lines show a downward trend, with the orange line generally having a higher loss than the blue line.

Epoch	Orange Line Model Loss	Blue Line Model Loss
0	0.99	0.99
1	0.95	0.90
2	0.85	0.75
3	0.75	0.65

epoch	Loss (LSTM)	Loss (LSTM+Attention)
0.0	0.95	0.65
1.0	0.52	0.35
2.0	0.32	0.22
3.0	0.25	0.18
4.0	0.20	0.15
5.0	0.15	0.12
6.0	0.12	0.10
7.0	0.10	0.08
8.0	0.08	0.07
9.0	0.07	0.06
10.0	0.06	0.06
11.0	0.05	0.05
12.0	0.04	0.05
13.0	0.04	0.05
14.0	0.03	0.05
15.0	0.03	0.05
16.0	0.02	0.06
17.0	0.02	0.05
18.0	0.02	0.05

BiLSTM with Attention

Let us create a BiLSTM model with Attention layer and the embeddings from FastText.

```
lstm_units = 128
dense = 8
w2vec_bilstm_attention_model = dl_pipeline.create_BiLSTM_Attention_nn_model(lstm_units, embedding_matrix, vocab)
```

```
w2vec_bilstm_attention_model.summary()
```

```
Model: "model"
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 100)	0
embedding (Embedding)	(None, 100, 300)	219200
bidirectional (Bidirectional)	(None, 100, 256)	439296
attention (attention)	(None, 256)	356
dense (Dense)	(None, 8)	2056

```
Total params: 670,908  
Trainable params: 441,708  
Non-trainable params: 229,200
```

Train model

I have experimented with different hyperparameters:

- epochs *batch_size
- Number of hidden units

```
epochs = 20  
batch_size = 128  
history = dl_pipeline.train_evaluate_model(w2vec_bilstm_attention_model, X_train, y_train, X_test,  
                                           y_test, epochs, batch_size)
```

Evaluate Model

Evaluate model with different metrics:

```

• Accuracy
• Precision
• Recall
• F1-score *AUC ROC Score

```

```

d1_pipeline.evaluate_model(w2vec_bilstm_attention_model, X_test, y_test)

```

```

13/13 [=====] - 4s 294ms/step

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	33
1	0.80	0.89	0.84	9
2	1.00	0.98	0.99	48
3	0.95	0.97	0.96	38
4	0.99	1.00	1.00	632
5	0.00	0.00	0.00	1
6	1.00	1.00	1.00	36
7	0.00	0.00	0.00	3
accuracy			0.99	800
macro avg	0.72	0.73	0.72	800
weighted avg	0.99	0.99	0.99	800

```

/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision
and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' pa

```

```

parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```

	0	1	2	3	4	5	6	7
0	33	0	0	0	0	0	0	0
1	0	8	0	0	1	0	0	0
2	0	0	47	0	1	0	0	0
3	0	0	0	37	1	0	0	0
4	0	0	0	1	60	0	0	1
5	0	0	0	0	0	1	0	0
6	0	0	0	0	0	0	36	0
7	0	2	0	1	0	0	0	0

Plot Accuracy and Loss

```
dl_pipeline.plot_accuracy_loss(history)
```

The plot displays the training progress of a deep learning pipeline. The blue line represents the model's accuracy, which increases steadily from about 0.945 at epoch 0 to a peak of approximately 0.995 at epoch 95, then slightly decreases. The orange line represents the model's loss, which also increases from about 0.945 at epoch 0 to a peak of approximately 0.995 at epoch 95, then slightly decreases. The legend indicates that the blue line represents 'model accuracy'.

Epoch	Accuracy (Blue Line)	Loss (Orange Line)
0	0.945	0.945
10	0.955	0.965
20	0.965	0.975
30	0.975	0.985
40	0.985	0.990
50	0.988	0.992
60	0.990	0.993
70	0.992	0.994
80	0.993	0.994
90	0.994	0.994
100	0.995	0.995

The figure consists of two line plots. The top plot shows training loss (blue line) and validation loss (orange line) over 18 epochs. The training loss starts at approximately 0.92 and decreases to about 0.94. The validation loss starts at approximately 0.93 and decreases to about 0.94. The bottom plot shows training loss (blue line) and validation loss (orange line) over 18 epochs. The training loss starts at approximately 1.25 and increases to about 1.25. The validation loss starts at approximately 0.92 and increases to about 1.05.

epoch	bert_encoder	bert_decoder	KerasLayer
0.0	0.45	0.45	0.45
2.5	0.35	0.35	0.35
5.0	0.25	0.25	0.25
7.5	0.22	0.22	0.22
10.0	0.20	0.20	0.20
12.5	0.18	0.18	0.18
15.0	0.17	0.17	0.17
17.5	0.16	0.16	0.16
20.0	0.18	0.18	0.18

BERT Model

Let us create a model with pretrained BERT embeddings. I will use the model provided by TensorFlow with KerasLayer wrapper. Let us define the preprocessor and embedding layers.

```
bert_preprocess = hub.KerasLayer("https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3")
bert_encoder = hub.KerasLayer("https://tfhub.dev/tensorflow/bert_en_uncased_L-12_H-768_A-12/4")
```

2022-07-09 20:21:14.318 : INFO : Using /tmp/tfhub_modules to cache modules.

```

2022-07-09 20:21:14,327: INFO: Downloaded TF-Rub Module 'https://tfhub.dev/tensorflow/bert_en_unscaled_preprocess/3'
2022-07-09 20:21:15,642: INFO: Downloaded https://tfhub.dev/tensorflow/bert_en_unscaled_preprocess/3. Total size: 1.96MB
2022-07-09 20:21:15,648: INFO: Downloaded TF-Rub Module 'https://tfhub.dev/tensorflow/bert_en_unscaled_preprocess/3'
2022-07-09 20:21:19,093: INFO: Downloading TF-Rub Module 'https://tfhub.dev/tensorflow/bert_en_unscaled_12_H-768_A-12/4'.
2022-07-09 20:21:35,954: INFO: Downloaded https://tfhub.dev/tensorflow/bert_en_unscaled_12_H-768_A-12/4. Total size: 429.6MB
2022-07-09 20:21:35,957: INFO: Downloaded TF-Rub Module 'https://tfhub.dev/tensorflow/bert_en_unscaled_12_H-768_A-12/4'.

dense = 8
bert_model = d1.pipeline.create_bert_model(bert_preprocess, bert_encoder, dense)

bert_model.summary()

Model: "model_1"

Layer (type)                 Output Shape                  Param #   Connected to
=====
text (InputLayer)            (None,)                      0         []
keras_layer (KerasLayer)     ('input_mask',) (None, 0) ('text[0][0]')

```

```
keras_layer_0 (DenseLayer) (None, 128)
    'input_type_ids':
      (None, 128),
      'input_word_ids':
      (None, 128)

keras_layer_1 (KerasLayer) (None, 128, 768)
    {'default': (None, 109482241), 'keras_layer[0][0]',
    768},
    'encoder_outputs':
    [(None, 128, 768),
    (None, 128, 768),
    (None, 128, 768),
    (None, 128, 768),
    (None, 128, 768),
    (None, 128, 768),
    (None, 128, 768),
    (None, 128, 768),
    (None, 128, 768),
    (None, 128, 768),
    (None, 128, 768),
    (None, 128, 768)],
    'pooled_output': {
    None, 768},
    'sequence_output':
    (None, 128, 768)

dropout (Dropout) (None, 768) 0
    [keras_layer_1][0][131]*
```

```

dropout (Dropout): 0.1000 (None, 8) 6152 ['dropout_0[0][0]']
output (Dense): 0.1000 (None, 8) 6152
Total params: 109,488,393
Trainable params: 6,152
Non-trainable params: 109,482,241

# V train and test for BERT model
X_train = np.array(train_df['text'])
X_test = np.array(test_df['text'])

Train model

I have experimented with different hyperparameters:



- epochs
- batch_size
- Number of hidden units



epochs = 20
batch size = 8

```

```
history = dl.pipeline.train_evaluate_model(hcrt_model, X_train, y_train, X_test,
                                          test_val_epochs, batch_size)
```

```
Epoch 1/20  
484/484 [====] 0.7305 - val_loss: 0.8141 - val_accuracy: 0.8462 - val_precision: 0.8356 - val_recall: 0.7094  
Epoch 2/20  
484/484 [====] 0.7499 - val_loss: 0.7751 - val_accuracy: 0.8501 - val_precision: 0.8867 - val_recall: 0.6923  
Epoch 3/20  
484/484 [====] 0.7647 - val_loss: 0.5797 - val_accuracy: 0.8558 - val_precision: 0.8507 - val_recall: 0.7839  
Epoch 4/20  
484/484 [====] 0.7799 - val_loss: 0.5333 - val_accuracy: 0.8582 - val_precision: 0.8570 - val_recall: 0.7994  
Epoch 5/20  
484/484 [====] 0.7921 - val_loss: 0.5256 - val_accuracy: 0.8666 - val_precision: 0.8575 - val_recall: 0.8153  
Epoch 6/20  
484/484 [====] 0.8022 - val_loss: 0.5252 - val_accuracy: 0.8954 - val_precision: 0.8330 - val_recall: 0.8061  
Epoch 7/20  
484/484 [====] 0.8092 - val_loss: 0.5130 - val_accuracy: 0.8981 - val_precision: 0.8583 - val_recall: 0.8035  
Epoch 8/20  
484/484 [====] 0.8143 - val_loss: 0.4816 - val_accuracy: 0.8981 - val_precision: 0.8496 - val_recall: 0.8628 - val_accuracy: 0.8791
```

```
Epoch 9/20
484/484: 0.8221 - val_loss: 0.4249 - val_accuracy: 0.8664 - val_precision: 0.8858 - val_recall: 0.8872
Epoch 10/20
484/484: 0.8249 - val_loss: 0.4258 - val_accuracy: 0.8652 - val_precision: 0.8769 - val_recall: 0.8397
Epoch 11/20
484/484: 0.8397 - val_loss: 0.4431 - val_accuracy: 0.8713 - val_precision: 0.9101 - val_recall: 0.8459
Epoch 12/20
484/484: 0.8430 - val_loss: 0.4111 - val_accuracy: 0.8666 - val_precision: 0.8866 - val_recall: 0.8407
Epoch 13/20
484/484: 0.8740 - val_loss: 0.4555 - val_accuracy: 0.8639 - val_precision: 0.8990 - val_recall: 0.8563
Epoch 14/20
484/484: 0.8807 - val_loss: 0.4058 - val_accuracy: 0.8729 - val_precision: 0.9267 - val_recall: 0.8501
Epoch 15/20
484/484: 0.8830 - val_loss: 0.4542 - val_accuracy: 0.8299 - val_accuracy: 0.9674 - precision: 0.8907
Epoch 16/20
484/484: 0.8534 - val_loss: 0.3900 - val_accuracy: 0.8695 - val_precision: 0.8892 - val_recall: 0.8563
Epoch 17/20
484/484: 0.8547 - val_loss: 0.3732 - val_accuracy: 0.8621 - val_accuracy: 0.9702 - precision: 0.9018
```

```
Epoch 18/20
484/484 [=====] - 50s 104ms/step - loss: 0.4049 - accuracy: 0.9697 - precision: 0.9007
recall: 0.9313 - val_loss: 0.4002 - val_accuracy: 0.9691 - val_precision: 0.8897 - val_recall: 0.8934
Epoch 19/20
484/484 [=====] - 51s 105ms/step - loss: 0.3941 - accuracy: 0.9701 - precision: 0.8978
recall: 0.8940 - val_loss: 0.3831 - val_accuracy: 0.9718 - val_precision: 0.8988 - val_recall: 0.8728
Epoch 20/20
484/484 [=====] - 50s 104ms/step - loss: 0.3860 - accuracy: 0.9712 - precision: 0.9037
recall: 0.8611 - val_loss: 0.3492 - val_accuracy: 0.9716 - val_precision: 0.9098 - val_recall: 0.7959
25/25 [=====] - 8s 333ms/step - loss: 0.2870 - accuracy: 0.9803 - precision: 0.9365 - recall: 0.9038
Test Accuracy, Test Precision, Test Recall: (0.2869974672794342, 0.9803125262240437, 0.9365285053894
0.059, 0.9037500228441858)
```

```

[1/13] /usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision
on and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' pa
rameter to control this behavior.
  warn_pr(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision
on and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' pa
rameter to control this behavior.
  warn_pr(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precisi
on and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' pa
rameter to control this behavior.
  warn_pr(average, modifier, msg_start, len(result))
precision    recall  f1-score   support

0         0.85         1.00         0.92         33
1         1.00         0.11         0.20         10
2         0.90         0.75         0.82         48
3         0.94         0.45         0.61         38
4         0.93         0.39         0.56         632
5         0.00         0.00         0.00         1
6         1.00         0.60         0.88         36
7         0.00         0.00         0.00         3

accuracy          0.70         0.61         0.65         800

```

Confusion Matrix

	0	1	2	3	4	5	6	7
0	33	0	0	0	0	0	0	0
1	1	1	0	0	7	0	0	0
2	0	0	36	0	12	0	0	0
3	5	0	0	17	16	0	0	0
4	0	0	4	0	626	0	0	2
5	0	0	0	0	1	0	0	0
6	0	0	0	0	0	8	0	28
7	0	0	0	1	2	0	0	0

Actual label

Predicted label

Plot Accuracy and Loss

```
d1_pipeline.plot_accuracy_loss(history)
```

The graph displays the accuracy of two models over 100 iterations. The orange line represents a model that consistently achieves higher accuracy than the blue line. Both models show an overall upward trend, with some fluctuations. The orange line starts at approximately 0.947 and ends at 0.974. The blue line starts at approximately 0.943 and ends at 0.972.

Iteration	Orange Line Accuracy	Blue Line Accuracy
0	0.947	0.943
10	0.955	0.950
20	0.958	0.955
30	0.967	0.960
40	0.956	0.962
50	0.966	0.964
60	0.972	0.968
70	0.967	0.966
80	0.974	0.968
90	0.969	0.969
100	0.974	0.972

The graph displays the training loss for two models over 18 epochs. The x-axis is labeled 'epoch' and ranges from 0.0 to 17.5. The y-axis is labeled 'loss' and ranges from 0.5 to 0.9. The blue line, representing the proposed method, starts at a loss of approximately 0.88 at epoch 0 and decreases to about 0.42 by epoch 18. The orange line, representing the baseline method, starts at a loss of approximately 0.82 at epoch 0 and decreases to about 0.40 by epoch 18. A shaded blue area around the blue line indicates the variance or confidence interval of the proposed method's performance.

epoch	proposed method (loss)	baseline method (loss)
0.0	0.88	0.82
1.0	0.72	0.78
2.0	0.65	0.58
3.0	0.62	0.53
4.0	0.56	0.52
5.0	0.53	0.52
6.0	0.52	0.52
7.0	0.50	0.48
8.0	0.48	0.43
9.0	0.47	0.40
10.0	0.46	0.40
11.0	0.45	0.42
12.0	0.44	0.40
13.0	0.44	0.38
14.0	0.43	0.38
15.0	0.43	0.38
16.0	0.42	0.38
17.0	0.42	0.38
18.0	0.42	0.40

Model Performance Comparison

These are the model performance metrics for different models with different embedding vectors. As the classification problem is imbalanced, we should also look at metrics like Precision, Recall etc, apart from accuracy.

	Accuracy	Weighted Avg Precision	Weighted Avg Recall	Weighted Avg F1-Score
BILSTM (Glove)	0.99	0.99	0.99	0.99
BILSTM with Attention (Glove)	0.99	0.99	0.99	0.99
BILSTM (Word2Vec)	0.99	1.0	0.99	0.99
BILSTM with Attention (Word2Vec)	0.99	0.99	0.99	0.99
BERT	0.93	0.93	0.93	0.92

- All models perform almost similarly across different metrics
- BILSTM with Word2Vec has the best performance, 1.0 weighted average precision
- It correctly classifies majority of instances from classes with low support as well
- BERT is computationally expensive, and requires a GPU for faster processing. It also has a lower performance

Improving model performance with data:

Model performance can be improved by increasing the amount of data. Deep learning models are known to perform better with a large amount of data, and it will also prevent overfitting.

The dataset has class imbalance, with 3 classes having 1% data. To improve performance, it will be beneficial to collect more data from these classes and create a balanced distribution across classes.

Model performance also depends on quality of the data, so it is important to collect data that is as diverse and varied as possible.

Data augmentation techniques to increase data can also be used. Some data augmentation techniques include

- Back Translation - Translating to another language and back, which gives a newer sentence with similar context.
- Random word insertion - Replace random words with synonyms.
- Random Deletion - Randomly delete some word from sentence.

These techniques can be used if it is difficult to obtain more data, the augmented data will be new for the model, but easier to generate and will help in generalisation.